



Phase 6.2 — Writing MF4 Files

1. Purpose of This Step

Phase 6.2 defines how the Linux side will **store real-time CAN frames into MF4 files** for later replay, analysis, machine learning, or predictive maintenance.

This step is necessary because:

- Predictive rules will need historical windows (5–300 seconds)
 - Fault trending requires continuity across power cycles
 - Cloud/off-board analytics may consume MF4 natively
 - ISO timestamp formats must be preserved for traceability
 - MF4 allows multi-signal expansion later (temperature, GPS, SoC, etc.)
-

2. What We Already Have

From Phase 5, we now have:

- ✓ Realtime CAN decode using DBC
- ✓ Stable MCU → Bridge → Python pipeline
- ✓ Continuous signal extraction (ECU_A, ECU_B, DCDC voltages)
- ✓ Validation with simulated degradation scenario

We are now adding:

- **Persistent storage**
 - **Timestamp correctness**
 - **Replay capability**
 - **Minimal schema for predictive analytics**
-

3. MF4 Requirements (For Our Use Case)

We are not implementing every ASAM feature — just what this product needs.

Functional requirements

Requirement	Status
Store raw CAN frames	Yes
Store decoded signals	Yes
Store timestamps	Yes (Linux time)
Store bus metadata	Optional
Multi-channel	Optional (future)
Incremental write	Yes
Replay compatible	Yes
Compression	Optional
Cloud uploadable	Yes (future Phase)

Non-functional requirements

Property	Target
Write latency	< 10ms
CPU load	Low
File size	Efficient
Logs rollover	Daily / Size-based
Durability	Power-loss resilient
Open tooling	Cantools / asammfdf / Vector

4. MF4 Contents (What Goes In)

For Phase 6, each MF4 file contains:

Raw channel:

`timestamp, can_id, dlc, data[8], flags`

Decoded channel (optional for now):

`timestamp, ECUA_Supply_Voltage, ECUB_Supply_Voltage,
DCDC_Output_Voltage`

Metadata:

`DBC version
App build version
ECU → harness mapping
Start/end timestamps`

5. Writing Strategy

Three writing strategies were considered:

A) Frame-by-frame write

- Lowest RAM usage
 - High overhead
 - Slow
- Rejected**

B) Buffered block write (preferred)

- Good latency
 - Efficient file structure
 - Compatible with replay
- Selected**

C) Full capture → write at end

- Requires large RAM
 - Risky on power loss
- Rejected**

Selected block size: 100–500 ms window

6. File Lifecycle Model

File rotation options:

Mode	Use case
Daily	Fleet logging
Hourly	Telematics
Size (e.g., 10MB)	Storage constrained
Event-based	DTC / predictive events

For demo: **Size-based (5MB) + Manual flush**

7. Interaction With Predictive Maintenance Logic

The pipeline becomes:

`CAN frames → DBC decode → Rules → MF4 write → (Cloud optional)`

It allows:

- retrospective debugging
- model training
- event validation
- regulatory compliance (in real products)

The predictive logic may:

- ✓ run online (low latency)
 - ✓ validate offline (historical signals)
-

8. How ASAM MDF4 Fits Telematics Systems

MDF4 solves 3 problems CSV/logs do not:

1. **Time alignment**
2. **Binary efficiency**
3. **Multi-channel merging**

This matches commercial telematics (OEMs, fleets, F1 motorsport, etc.)

9. Implementation Plan (Phase 6.3 Preview)

Linux implementation will include:

- `MdfWriter` class
 - circular buffer for CAN frames
 - rollover logic
 - replay CLI/tooling
 - test harness using synthetic degradation signals
-

10. Validation for Phase 6.2

Success criteria:

- ✓ MF4 file is created
- ✓ Timestamps are correct
- ✓ Data replay matches original inputs
- ✓ File can be opened with `asammdf` or Vector tools

Stretch goals:

- decoded channel export
 - event-based file tagging
-

11. Risks & Notes

Risk	Mitigation
File corruption on power loss	journaled writes
CPU overhead	buffered writes
DBC version mismatch	metadata versioning
Flash wear	rotation strategy

12. Next Step

Proceed to **6.3 — Data replay & verification**

Where we will:

- replay MF4 → Bridge → decoder (same pipeline)
- confirm round-trip correctness
- prep the system for Phase 7 predictive rules